

## 1 Housekeeping

v0.0

First make sure HBase is running, with the following command in \$HBASE\_HOME: `$ ./bin/start-hbase.sh`

The HBase Shell is a Ruby script in `$(HBASE_HOME)/bin/hbase.rb` and to start the shell use: `$ hbase shell`

You can get an overview of all tables present in HBase using the `list` command (or restrict with `list 'test.*'`).

To get a grip of how a certain table is structured use the `describe` command: `describe 'test_table'`

To check how the cluster is doing use: `status 'simple' | 'summary' | 'detailed'` (note that these are three alternative parameters, providing different level of details; only specify one of them).

If you need to shut down the cluster use `shutdown` and to exit the shell use `exit`

## 2 Data Definition

**create** ... to create a table. Pass table name, a dictionary of specifications per column family, and optionally a dictionary of table configuration.

*Example* `create 'test_table', 'cf1', 'cf2'`

**alter** ... to alter the column family (CF) schema of a table. Note: the table must be disabled.

*Examples*

To add the 'cf1' column family in table 't1':

`alter 't1', NAME => 'f1'`

To delete the 'cf2' column family from table 't2':

`alter 't2', {NAME => 'cf2', METHOD => 'delete'}`

**drop** ... to drop a table. Note: the table must be disabled.

**disable** ... to disable a table, for example to drop it or to alter the schema.

**enable** ... to enable a table after altering the CF schema.

## 3 Data Manipulation

**put** ... to put a 'value' into at specified table/row/column combination, also known as a 'cell' in HBase.

*Example* `put 't1', 'row-key-1', 'cf1:col2', 'value'`

**delete** ... to mark a cell value as deleted. Note: deletes must match the deleted cell's coordinates (table/row/column) exactly. When scanning, a delete cell suppresses older versions.

*Example* `delete 't1', 'row-key-1', 'cf1:col2'`

## 4 Query

**count** ... to count the number of rows in a table. Note: this operation may take a long time.

*Example* `count 'my_table'`

**get** ... to retrieve a single cell.

*Examples*

To retrieve a cell with a certain row key:

`get 't1', 'row-`

`key-1'`

To retrieve a certain CF of a cell with a certain row key:

`get 't1', 'row-`

`key-1', COLUMN => 'cf1'`

To retrieve a certain version of a cell with a certain row key:

`get 't1', 'row-`

`key-1', TIMESTAMP =>`


`1371113904834`

**scan** ... to scan a table. Scanner specifications may include one or more of the following: `FILTER`, `COLUMNS`, `LIMIT`, `STARTROW`, `STOPROW`, or `TIMESTAMP`. If no `COLUMNS` are specified, all columns will be scanned. To scan all columns of a column family, leave the qualifier empty as in `'col_family:'`

*Examples*

See the next page, in the walk-through example.

- <http://wiki.apache.org/hadoop/Hbase/Shell>
- <http://refcardz.dzone.com/refcardz/hbase>
- <http://hbase.apache.org/book/book.html>
- <https://github.com/larsgeorge/hbase-book>



## 5 Walk-through Example

v0.0

```
$ hbase shell

hbase> create 'orders', 'client', 'product'

hbase> describe 'orders'

hbase> put 'orders', 'joe_2013-01-13', 'client:name', 'Joe'

hbase> put 'orders', 'joe_2013-01-13', 'client:address', 'Hillroad 1, SF'

hbase> put 'orders', 'joe_2013-01-13', 'product:title', 'iPhone 5'

hbase> put 'orders', 'joe_2013-01-13', 'product:delivery', '2013-01-13'

hbase> scan 'orders'

hbase> put 'orders', 'jane_2013-02-05', 'client:name', 'Jane'

hbase> put 'orders', 'jane_2013-02-05', 'client:address', 'Sunset Drive 42, NY'

hbase> put 'orders', 'jane_2013-02-05', 'product:title', 'Samsung S4'

hbase> put 'orders', 'jane_2013-02-05', 'product:delivery', '2013-05-02'

hbase> get 'orders', 'jane_2013-02-05', {COLUMN => 'product:'}

hbase> scan 'orders', FILTER => "ValueFilter(=,'substring:Sunset')"

hbase> scan 'orders', { COLUMNS => ['client'], FILTER =>
    "ValueFilter(=,'substring:road')" }

hbase> scan 'orders', { COLUMNS => ['product'], FILTER =>
    "ValueFilter(=,'substring:2013-') " }

hbase> disable 'orders'

hbase> drop 'orders'

hbase> list

hbase> exit
```

Note: This walk-through example creates a table `orders`, adds some data (via `put`) and then retrieves the data using `get` and `scan`. Finally, the table is removed again, using `disable` and `drop`.

## 6 Other Useful Stuff

To execute HBase commands **scripting-style**, use:

```
$ {HBASE_HOME}/bin/hbase shell PATH_TO_SCRIPT
```

For a comprehensive reference on **filters** (to be used in `FILTER => "..."` within `scan`) see the Jira issue HBASE-4176 'Exposing HBase Filters to the Thrift API' at <https://issues.apache.org/jira/browse/HBASE-4176> which has a MS Word document attached: <https://issues.apache.org/jira/secure/attachment/12490608/Filter%20Language%283%29.docx>

Also, regarding filters, note that the following **compare operators** are available:

- LESS ... <
- LESS\_OR\_EQUAL ... <=
- EQUAL ... =
- NOT\_EQUAL ... !=
- GREATER\_OR\_EQUAL ... >=
- GREATER ... >
- NO\_OP